

Regular Expression (regex)

Jonathan Feinberg
Dept. of Informatics, Univ. of Oslo
Simula Research Laboratory

August 2014



Regular expression (regex)

Contents

- Motivation for regular expression
- Regular expression syntax
- Lots of examples on problem solving with regular expressions
- For more information on regex in Python:

`http://uio-inf3331.github.io/resources-14/
doc/texts/pub/pyregex.html`



Digression: Editor war!



Regular expression is advanced search and replace

It is strongly encourage to learn by experimenting in a supported editor.

- Emacs

`lmgify.com/?q=emacs+tutorial`

- Vim

```
$ sudo apt-get install vim-gnome # or vim-gtk or vim-common
$ vimtutor           # 30 minutes to learn basic stuff
```

- Sublime

`http://www.sublimetext.com/`

- Notepad++ (Windows)

`http://notepad-plus-plus.org/`

- textmate (Mac)

`http://macromates.com/`

Regex search and replace in Emacs and Vim

🟦 Emacs search:

C-M-S

```
M-x isearch-forward-regex<enter>
```

🟦 Emacs replace:

C-M-%

```
M-x replace-regexp<enter>
```

```
# [y]es/[n]o/all[!]/cancel[<C-g>]
```

🟦 Vim highlighting:

```
:set hls<enter>
```

```
:set hlsearch<enter>
```

🟦 Vim search:

$$\wedge \backslash v$$

- Vim replace:

$$: \%S / \backslash v$$

Show and tell

The command line wild card

- Consider a simulation code with this type of output:

```
$ ls *.tex  
report.tex
```

- Here '*' is a wild card and can be read as:
“Zero or more of any non-linefeed character type”.
- Simple and powerful, but some what unrefined.

Regular expression introduction

- Classical wild card split into two components:

- '.' : Any character

- '*' : Zero or more

- Replacate wild card with '.*':

```
$ grep -E ".*" unsorted_fruits
orange
pear
apple
grape
pineapple
```

```
$ grep -E ".*apple" unsorted_fruits
apple
pineapple
```

Refined search

Character		Quantifier	
.	Non-linefeed character	*	Zero or more
\w (\W)	(Not) alphabet character	+	One or more
\d (\D)	(Not) numerical character	?	Zero or one
\s (\S)	(Not) white space	{ 4 }	Custom quantifier
[A-F\d_]	Custom character	{ 2 , 4 }	Range quantifier
[^abc]	All but custom character	{ , 4 }	Open ended range

Only strings ending with 'apple':

```
$ grep -E "\w+apple" unsorted_fruits
pineapple
```

Delimiters

- A text have multiple “delimiters”, i.e. structures in text which are not characters:

```
      This      is      some
      ^         <  >      $ \n
      text      worth      reading.
      ^         \b      \b          $
```

- Find apples but not pineapples:

```
$ grep -E "\bapple\b" unsorted_fruits
apple
```

Inconsistent Regex

- Reserved characters that has to be “cancelled” with a ‘\’ prefix in Bash/Python:

^	[.	\$	{	*	(+)		?	<	>
\w	\W	\d	\D	\s	\S	\b	\B					

- Emacs’ isearch-forward-regex and replace-regex:

^	[.	\$	\{	*	\(+	\)	\	?	\<	\>
\w	\W			\s-	\S-	\b	\B					

- Vim’s “magic mode”:

^	\[.	\$	\{	*	\(\+	\)	\	\?	\<	\>
\w	\W	\d	\D	\s	\S	\b	\B					

- Vim’s “very magic mode”:

^	[.	\$	{	*	(+)		?	<	>
\w	\W	\d	\D	\s	\S	\b	\B					

Regular expression in Python

- Find all words:

```
>>> import re
>>> text = "apple, apples, pineapple and appletini"
>>> print re.findall(r"\w+", text)
['apple', 'apples', 'pineapple', 'and', 'appletini']
```

- Replace words ending with “apple” with “pear”:

```
>>> re.sub(r"apple\b", "pear", text)
'pear, apples, pinepear and appletini'
```

- As always, read more in the documentation:

```
$ pydoc re
```

Extraction

- Parenthesis can be used to extract sub-strings:

```
>>> re.findall(r"\bapple(\w*)", text)
['', 's', 'tini']
```

- Multiple extraction is possible:

```
>>> re.findall(r"(\w*)apple(\w*)", text)
[('', ''), ('', 's'), ('pine', ''), ('', 'tini')]
```

- Extraction can be used in advanced substitution:

```
>>> re.sub(r"(\w*)apple(\w*)", r"\2apple\1", text)
'apple, sapple, applepine and tiniapple'
```

\1, \2, ..., \9 represent extractions.

\0 is the full match.

Groupings

- Paranthesis can also be used to to group alternatives through the '|' character.

- For example, find words shorter than 6 charracters, but also longer than 7:

```
>>> re.findall(r"(\w{8,}|\w{1,5})", text)
['apple', 'pineapple', 'appletini']
```

- Note: Precedence order from left to right:

```
>>> re.findall(r"(\w{1,5}|\w{8,})", text)
['apple', 'apple', 's', 'pinea', 'pple', 'apple', 'tini']
```

Left clause must fail before right clause can be addressed.

Imprecise syntax leads to too much greed

- Finding all words starting with “a” and end with “s”:

```
>>> re.findall(r"a.*s", text)
['apple, apples']
```

Ups!

- Quantifiers like '*' and '+' are greedy.
- Quantifiers can be made non-greedy by placing a '?' after them:

```
>>> re.findall(r"a.*?s", text)
['apples']
```
- Note 1: '?' after characters are still “zero or one”.
- Note 2: '?' is not supported in Vim. Use negative range instead:
'{-0,}'

Larger example: extracting all the numbers!

```
t=2.5  a: 1.0 6.2 -2.2    12 iterations and eps=1.38756E-05
t=4.25 a: 1.0 1.4    6 iterations and eps=2.22433E-05
>> switching from method AQ4 to AQP1
t=5    a: 0.9    2 iterations and eps=3.78796E-05
t=6.386 a: 1.0 1.1525    6 iterations and eps=2.22433E-06
>> switching from method AQP1 to AQ2
t=8.05  a: 1.0    3 iterations and eps=9.11111E-04
...
```

- Different ways of writing real numbers:
-3, 42.9873, 1.23E+1, 1.2300E+01, 1.23e+01
- Three basic forms:
 - integer: -3
 - decimal notation: 42.9873, .376, 3.
 - scientific notation: 1.23E+1, 1.2300E+01, 1.23e+01, 1e1

A simple approach

- Could just collect the legal characters in the three notations:

`[0-9.Ee\-\+]+`

- Downside: this matches text like

12-24
24.-
--E1--
+++++

- How can we define precise regular expressions for the three notations?

Decimal notation regex

- Regex for decimal notation:

```
-?\d*\.\d+
```

```
# or in emacs:
```

```
-?[0-9]*\.[0-9]+
```

- Problem: this regex does not match '3.'

- The fix

```
-?\d*\.\d*
```

is ok but matches text like '-.' and (much worse!) ''

- Trying it on

```
'some text. 4. is a number.'
```

gives a match for the first period!

Fix of decimal notation regex

- We need a digit before OR after the dot

- The fix:

`-?(\d*\.\d+|\d+\.\d*)`

- A more compact version (just "OR-ing" numbers without digits after the dot):

`-?(\d*\.\d+|\d+\.)`

Combining regular expressions

- Make a regex for integer or decimal notation:

`(integer OR decimal notation)`

using the OR operator and parenthesis:

`-?(\d+|(\d+\.\d*|\d*\.\d+))`

- Problem: `22.432` gives a match for `22`
(i.e., just digits? yes - `22` - match!)

Check the order in combinations!

- Remedy: test for the most complicated pattern first

(decimal notation OR integer)

```
-?((\d+\.\d*|\d*\.\d+)|\d+)
```

- Modularize the regex:

```
real_in = r'\d+'
```

```
real_dn = r'(\d+\.\d*|\d*\.\d+).'
```

```
real = '-?(' + real_dn + '|' + real_in + ')'
```

Scientific notation regex (1)

- Write a regex for numbers in scientific notation

- Typical text: 1.27635E+01, -1.27635e+1

- Regular expression:

`-?\d\.\d+[Ee][+\-]\d\d?`

- = optional minus, one digit, dot, at least one digit, E or e, plus or minus, one digit, optional digit

Scientific notation regex (2)

- Problem: `1e+00` and `1e1` are not handled
- Remedy: zero or more digits behind the dot, optional `e/E`, optional sign in exponent, more digits in the exponent (`1e001`):

`-?\d\.\d*[Ee][+\-]?\d+`

Making the regex more compact

- A pattern for integer or decimal notation:

`-? ((\d+\.\d* | \d*\.\d+) | \d+)`

- Can get rid of an OR by allowing the dot and digits behind the dot be optional:

`-? (\d+ (\.\d*)? | \d*\.\d+)`

- Such a number, followed by an optional exponent (a la `e+02`), makes up a general real number (!)

`-? (\d+ (\.\d*)? | \d*\.\d+) ([eE] [+|-]? \d+)?`

A more readable regex

- Scientific OR decimal OR integer notation:

```
-?(\d\.\?\d*[Ee][+|-]?\d+|(\d+\.\d*|\d*\.\d+)|\d+)
```

or better (modularized):

```
real_in = r'\d+'
```

```
real_dn = r'(\d+\.\d*|\d*\.\d+).'
```

```
real_sn = r'(\d\.\?\d*[Ee][+|-]?\d+).'
```

```
real = '-?(' + real_sn + '|' + real_dn + '|' + real_in + ')'
```

Grab the groups

- Enclose parts of a regex in () to extract the parts:

```
pattern = r"t=(.*)\s+a:.*\s+(\d+)\s+.*=(.*)"  
# groups:      ( )              ( )              ( )
```

This defines three groups (t, iterations, eps)

- In Python code:

```
matches = re.findall(pattern, line)  
for match in matches:  
    time = float(match[0])  
    iter = int    (match[1])  
    eps  = float(match[2])
```

Pattern-matching modifiers (1)

- ...also called flags in Python regex documentation

- Check if a user has written "yes" as answer:

```
re.findall('yes', answer)
```

- Problem: "YES" is not recognized; try a fix

```
re.findall(r'(yes|YES)', answer)
```

- Should allow "Yes" and "YEs" too...

```
re.findall(r'[yY][eE][sS]', answer)
```

- This is hard to read and case-insensitive matches occur frequently - there must be a better way!

Pattern-matching modifiers (2)

```
matches = re.findall('yes', answer, re.IGNORECASE)
# pattern-matching modifier: re.IGNORECASE
# now we get a match for 'yes', 'YES', 'Yes' ...

# ignore case:
re.I or re.IGNORECASE

# let ^ and $ match at the beginning and
# end of every line:
re.M or re.MULTILINE

# allow comments and white space:
re.X or re.VERBOSE

# let . (dot) match newline too:
re.S or re.DOTALL

# let e.g. \w match special chars (?, ?, ...):
re.L or re.LOCALE
```

Comments in a regex

- The `re.X` or `re.VERBOSE` modifier is very useful for inserting comments explaining various parts of a regular expression

- Example:

```
# real number in scientific notation:
```

```
real_sn = r"""
```

```
-?          # optional minus
```

```
\d\.\d+      # a number like 1.4098
```

```
[Ee][+\-]\d\d? # exponent, E-03, e-3, E+12
```

```
"""
```

```
match = re.search(real_sn, 'text with a=1.92E-04 ',  
                  re.VERBOSE)
```

```
# or when using compile:
```

```
c = re.compile(real_sn, re.VERBOSE)
```

```
match = c.findall('text with a=1.9672E-04 ')
```

Substitution example

- Suppose you have written a C library which has many users
- One day you decide that the function

```
void superLibFunc(char* method, float x)
```

would be more natural to use if its arguments were swapped:

```
void superLibFunc(float x, char* method)
```

- All users of your library must then update their application codes - can you automate?

Substitution with backreferences

- You want locate all strings on the form

```
superLibFunc(arg1, arg2)
```

and transform them to

```
superLibFunc(arg2, arg1)
```

- Let `arg1` and `arg2` be groups in the regex for the `superLibFunc` calls

- Write out

```
superLibFunc(\2, \1)
```

```
# recall: \1 is group 1, \2 is group 2 in a re.sub command
```


Regex for the function calls (1)

- Basic structure of the regex of calls:

```
superLibFunc\s*\(\s*arg1\s*,\s*arg2\s*\)
```

but what should the `arg1` and `arg2` patterns look like?

- Natural start: `arg1` and `arg2` are valid C variable names

```
arg = r"[A-Za-z_0-9]+"
```

- Fix; digits are not allowed as the first character:

```
arg = "[A-Za-z_][A-Za-z_0-9]*"
```

Regex for the function calls (2)

• The regex

```
arg = "[A-Za-z_][A-Za-z_0-9]*"
```

works well for calls with variables, but we can call `superLibFunc` with numbers too:

```
superLibFunc ("relaxation", 1.432E-02);
```

• Possible fix:

```
arg = r"[A-Za-z0-9_.\-+\\" ]+"
```

but the disadvantage is that `arg` now also matches

```
.-+32skj 3.ejks
```

Constructing a precise regex (1)

- Since `arg2` is a float we can make a precise regex: legal C variable name OR legal real variable format

```
arg2 = r"([A-Za-z_][A-Za-z_0-9]*|" + real + \
        "|float\s+[A-Za-z_][A-Za-z_0-9]*" + ")"
```

where `real` is our regex for formatted real numbers:

```
real_in = r"-?\d+"
real_sn = r"-?\d\.\d+[Ee][+|-]\d\d?"
real_dn = r"-?\d*\.\d+"
real = r"\s*(" + real_sn + "|" + real_dn + "|" + real_in + r")\s+"
```

Constructing a precise regex (2)

- We can now treat variables and numbers in calls
- Another problem: should swap arguments in a user's definition of the function:

```
void superLibFunc(char* method, float x)
```

to

```
void superLibFunc(float x, char* method)
```

Note: the argument names (`x` and `method`) can also be omitted!

- Calls and declarations of `superLibFunc` can be written on more than one line and with embedded C comments!
- Giving up?

A simple regex may be sufficient

- Instead of trying to make a precise regex, let us make a very simple one:

```
arg = ' .+ '    # any text
```

- "Any text" may be precise enough since we have the surrounding structure,

```
superLibFunc\s*(\s*arg\s*,\s*arg\s*)
```

and assume that a C compiler has checked that `arg` is a valid C code text in this context

Refining the simple regex

- A problem with `.+` appears in lines with more than one calls:

```
superLibFunc(a,x);  superLibFunc(ppp,qqq);
```

- We get a match for the first argument equal to

```
a,x);  superLibFunc(ppp
```

- Remedy: non-greedy regex (see later) or

```
arg = r"[^,]+"
```

This one matches multi-line calls/declarations, also with embedded comments (`.+` does not match newline unless the `re.S` modifier is used)

Swapping of the arguments

● Central code statements:

```
arg = r"[^,]+"
call = r"superLibFunc\s*\(\s*(%s),\s*(%s)\)" % (arg,arg)

# load file into filestr

# substitute:
filestr = re.sub(call, r"superLibFunc(\2, \1)", filestr)

# write out file again
fileobject.write(filestr)
```

Testing the code

• Test text:

```
superLibFunc(a,x);  superLibFunc(qqq,ppp);
superLibFunc ( method1, method2 );
superLibFunc(3method /* illegal name! */, method2 ) ;
superLibFunc(  _method1,method_2) ;
superLibFunc (
    method1 /* the first method we have */ ,
    super_method4 /* a special method that
                        deserves a two-line comment... */
) ;
```

• The simple regex successfully transforms this into

```
superLibFunc(x, a);  superLibFunc(ppp, qqq);
superLibFunc(method2 , method1);
superLibFunc(method2 , 3method /* illegal name! */) ;
superLibFunc(method_2, _method1) ;
superLibFunc(super_method4 /* a special method that
                        deserves a two-line comment... */
, method1 /* the first method we have */ ) ;
```

• Notice how powerful a small regex can be!!

• Downside: cannot handle a function call as argument

Shortcomings

- The simple regex

`[^,]+`

breaks down for comments with comma(s) and function calls as arguments, e.g.,

```
superLibFunc(m1, a /* large, random number */);  
superLibFunc(m1, generate(c, q2));
```

The regex will match the longest possible string ending with a comma, in the first line

`m1, a /* large,`

but then there are no more commas ...

- A complete solution should *parse* the C code

More easy-to-read regex

- The superLibFunc call with comments and named groups:

```
call = re.compile(r"""
    superLibFunc    # name of function to match
    \s*             # possible whitespace
    \(              # parenthesis before argument list
    \s*             # possible whitespace
    (?P<arg1>%s)     # first argument plus optional whitespace
    ,               # comma between the arguments
    \s*             # possible whitespace
    (?P<arg2>%s)     # second argument plus optional whitespace
    \)              # closing parenthesis
    """ % (arg,arg), re.VERBOSE)

# the substitution command:
filestr = call.sub(r"superLibFunc(\g<arg2>,
                               \g<arg1>)",filestr)
```

Example

- Goal: remove C++/Java comments from source codes

- Load a source code file into a string:

```
filestr = open(somefile, 'r').read()  
# note: newlines are a part of filestr
```

- Substitute comments *// some text...* by an empty string:

```
filestr = re.sub(r'//.*', '', filestr)
```

- Note: . (dot) does not match newline; if it did, we would need to say

```
filestr = re.sub(r'//[^\n]*', '', filestr)
```

Failure of a simple regex

- How will the substitution

```
filestr = re.sub(r'//[^\n]*', '', filestr)
```

treat a line like

```
const char* heading = "-----//-----";
```

???